KubeCon | CloudNativeCon Europe 2019

Autoscaling Multi-Cluster Observability with Thanos and Linkerd

Frederic Branczyk @brancz @fredbrancz

Andrew Seigner @siggy

What is a service mesh?

Control plane

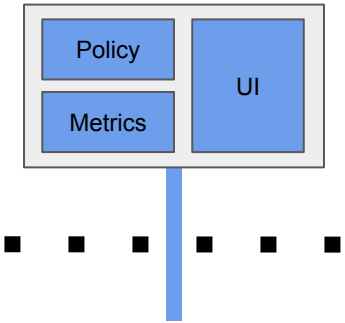

Data plane

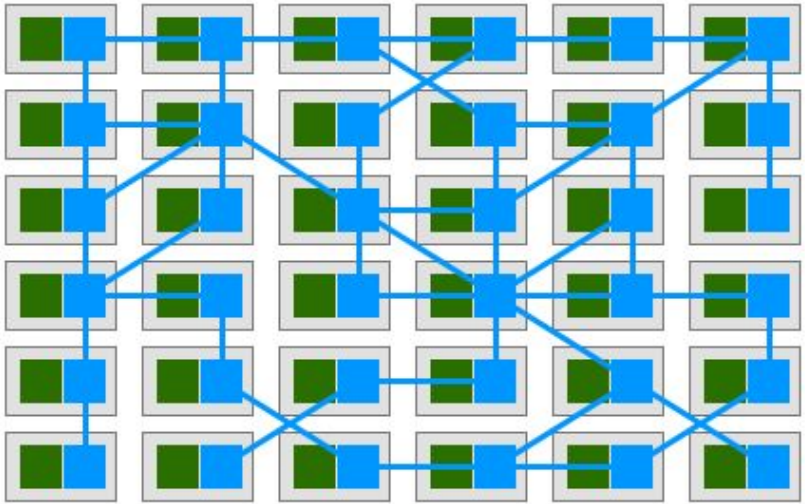

Linkerd 1.x

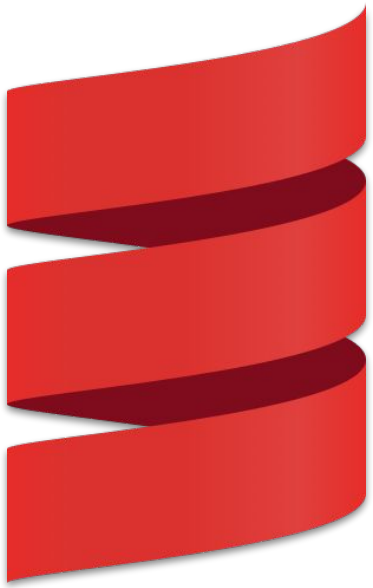

Linkerd 2.0

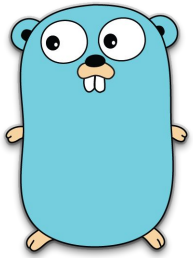

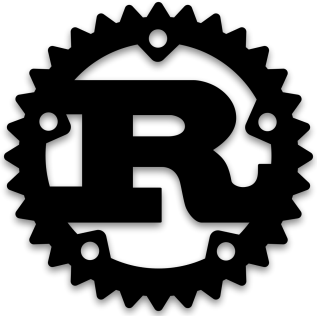

Linkerd 2.0

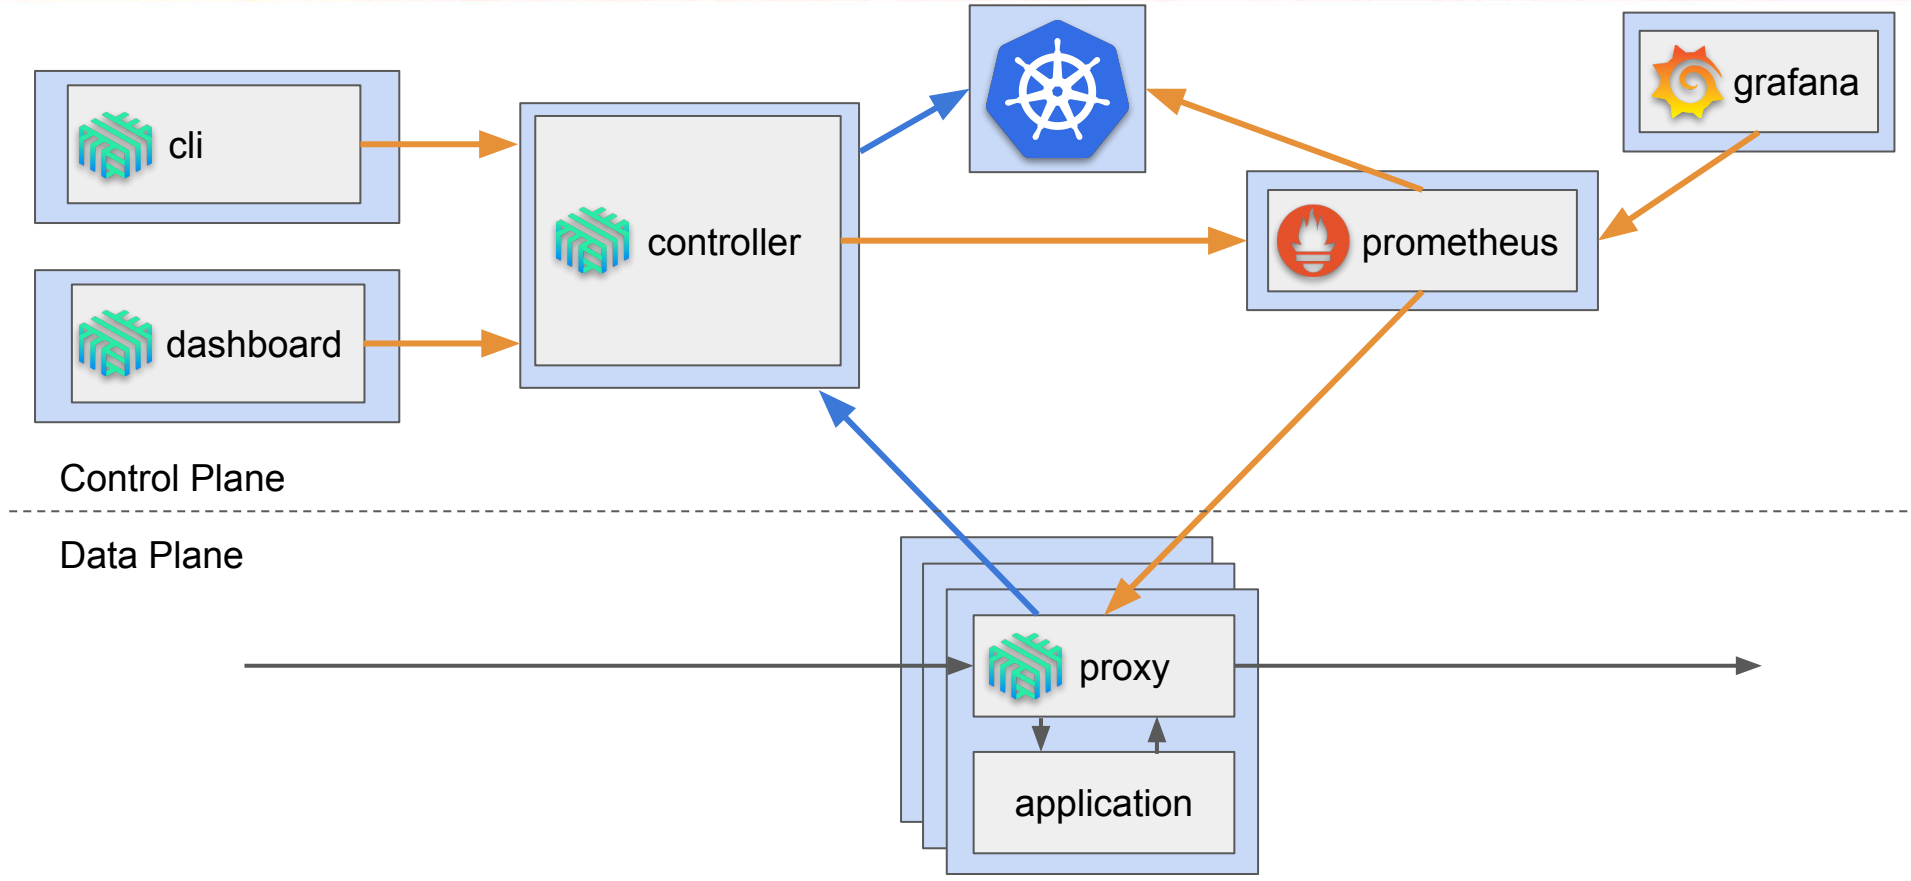
cli
dashboard
controller
grafana
prometheus
Control Plane
Data Plane
proxy
application

Linkerd Prometheus Endpoint

KubeCon CloudNativeCon Europe 2019

```
$ curl localhost:4191/metrics
request_total{tls="true",...} 123
response_total{tls="true",status_code="200",...} 123
response_latency_ms_bucket{tls="true",...} 120
```

Quick Prometheus intro

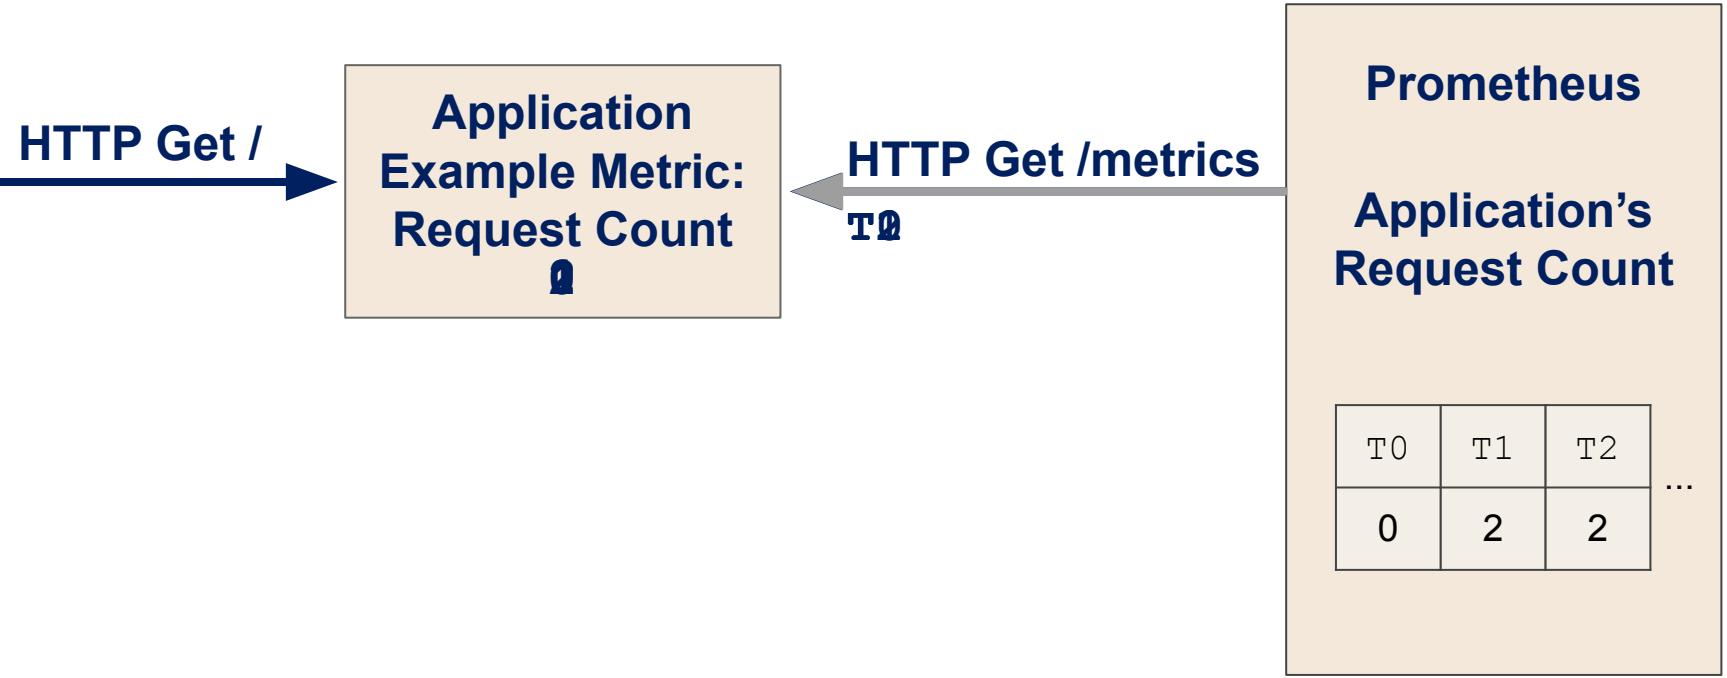

Typical monitoring setup

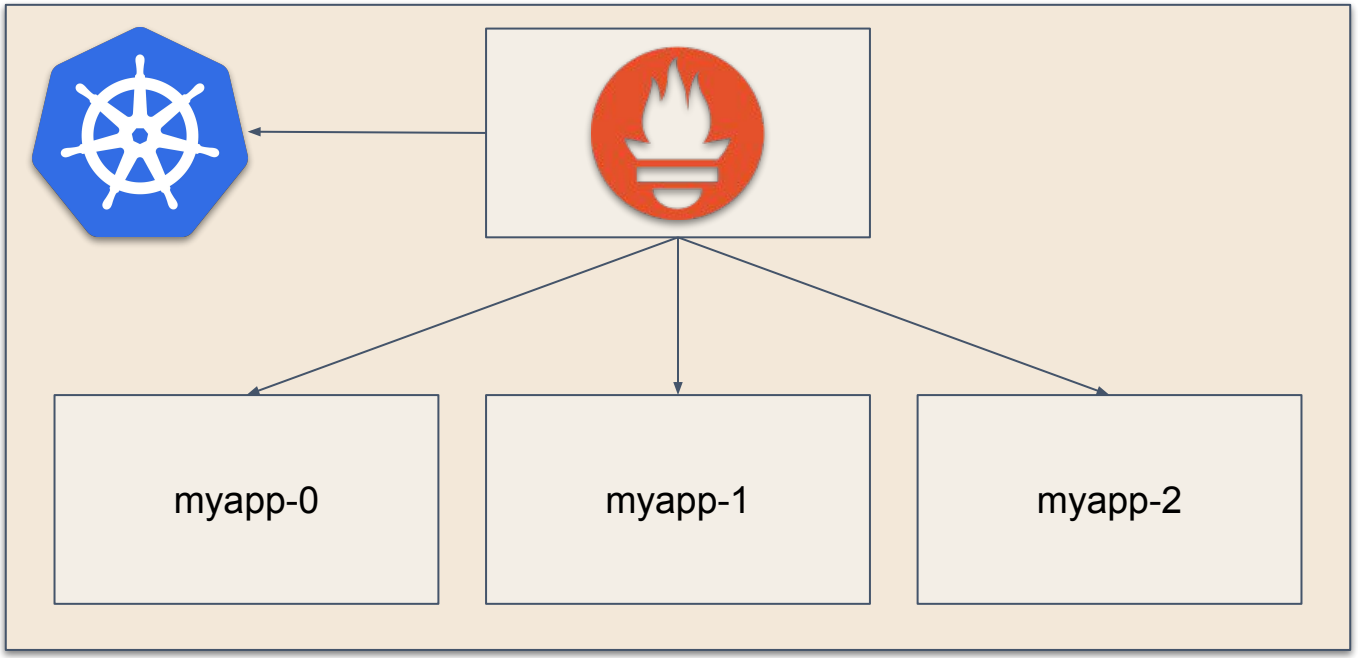

Most requested features?

- Long term storage
- Global view
- Downsampling

Thanos

- Long term storage
- Global view
- Downsampling
- Distributed Prometheus

Anatomy of Thanos

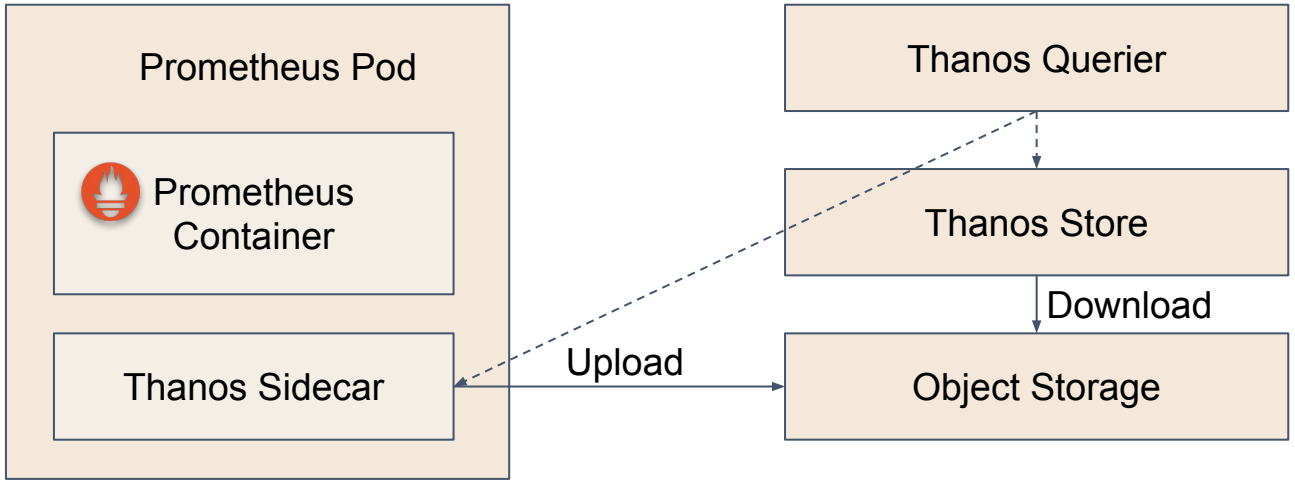

Thanos receive

- New component
- Accepts remote-write from Prometheus
- http://bit.ly/thanos-receive-proposal

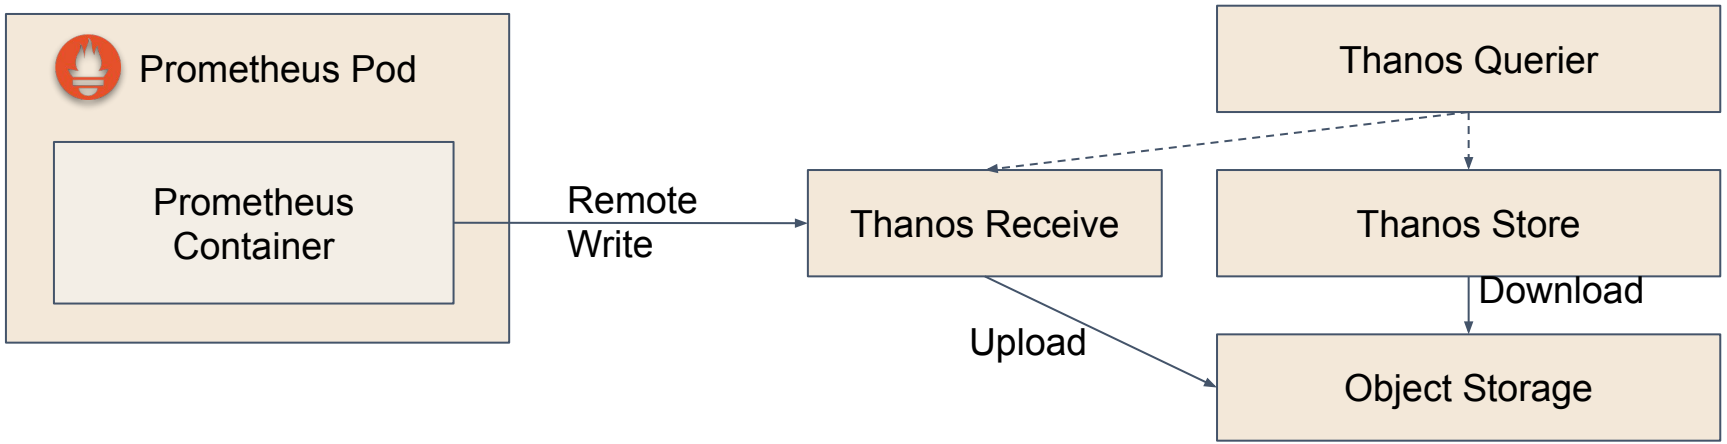

Architect to your need

So Thanos = Multi cluster?

StoreAPI

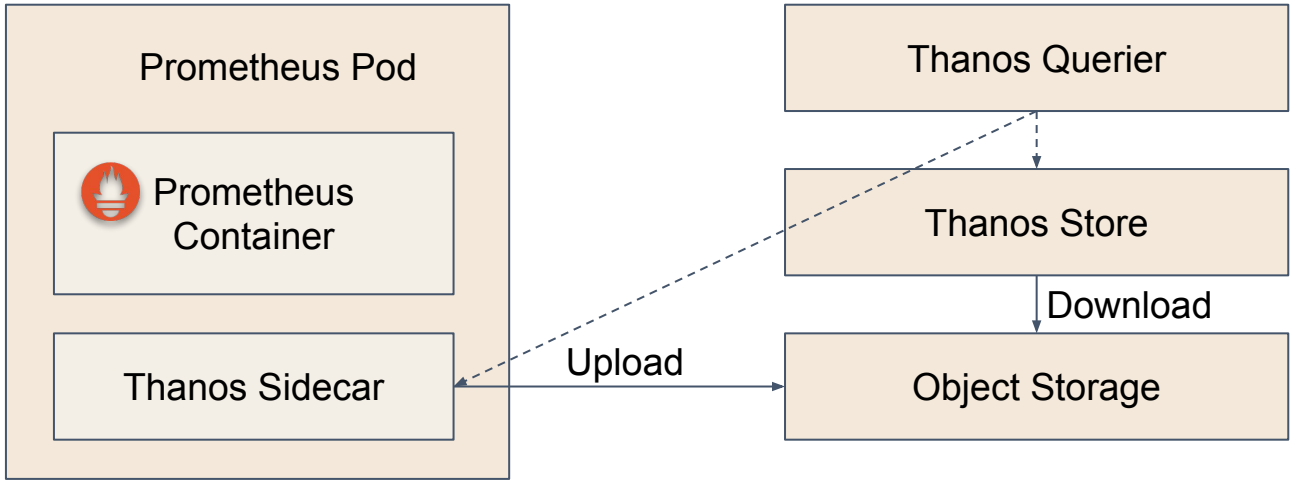

---------- StoreAPI

StoreAPI is a building block

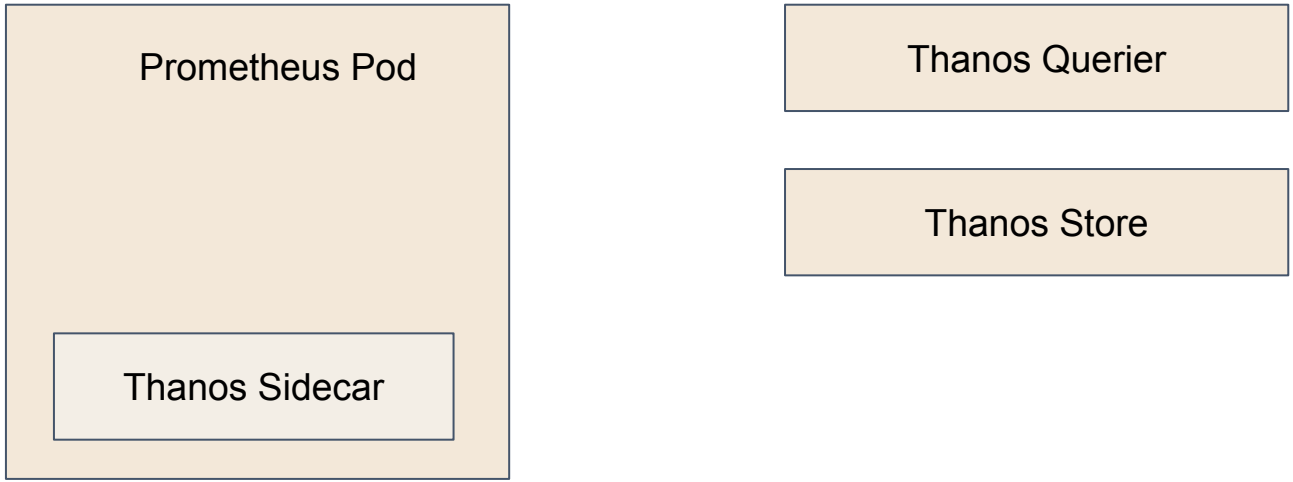

StoreAPI is a building block

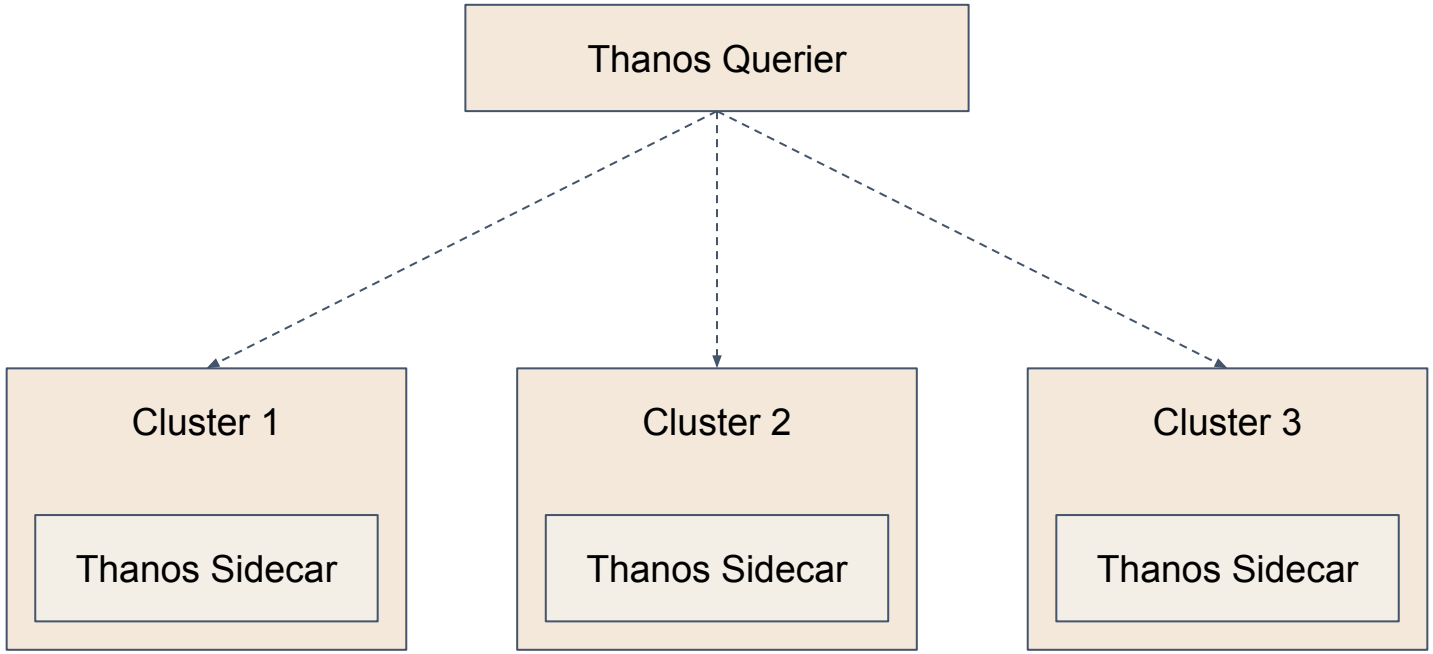

StoreAPI is a building block

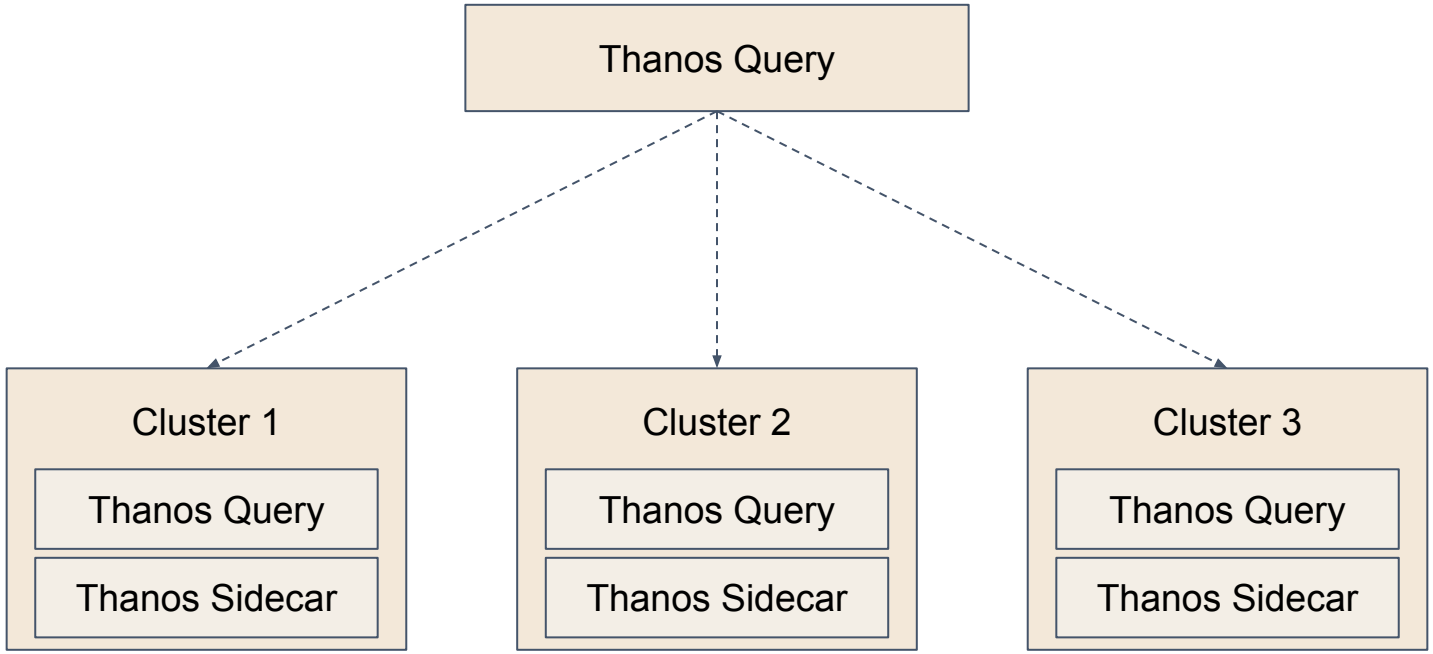

StoreAPI is a building block

- Thanos is about deployment flexibility
- StoreAPI provides this flexibility
  - Each component exposes it
  - Stackable

Demo time!

bit.ly/thanos-linkerd

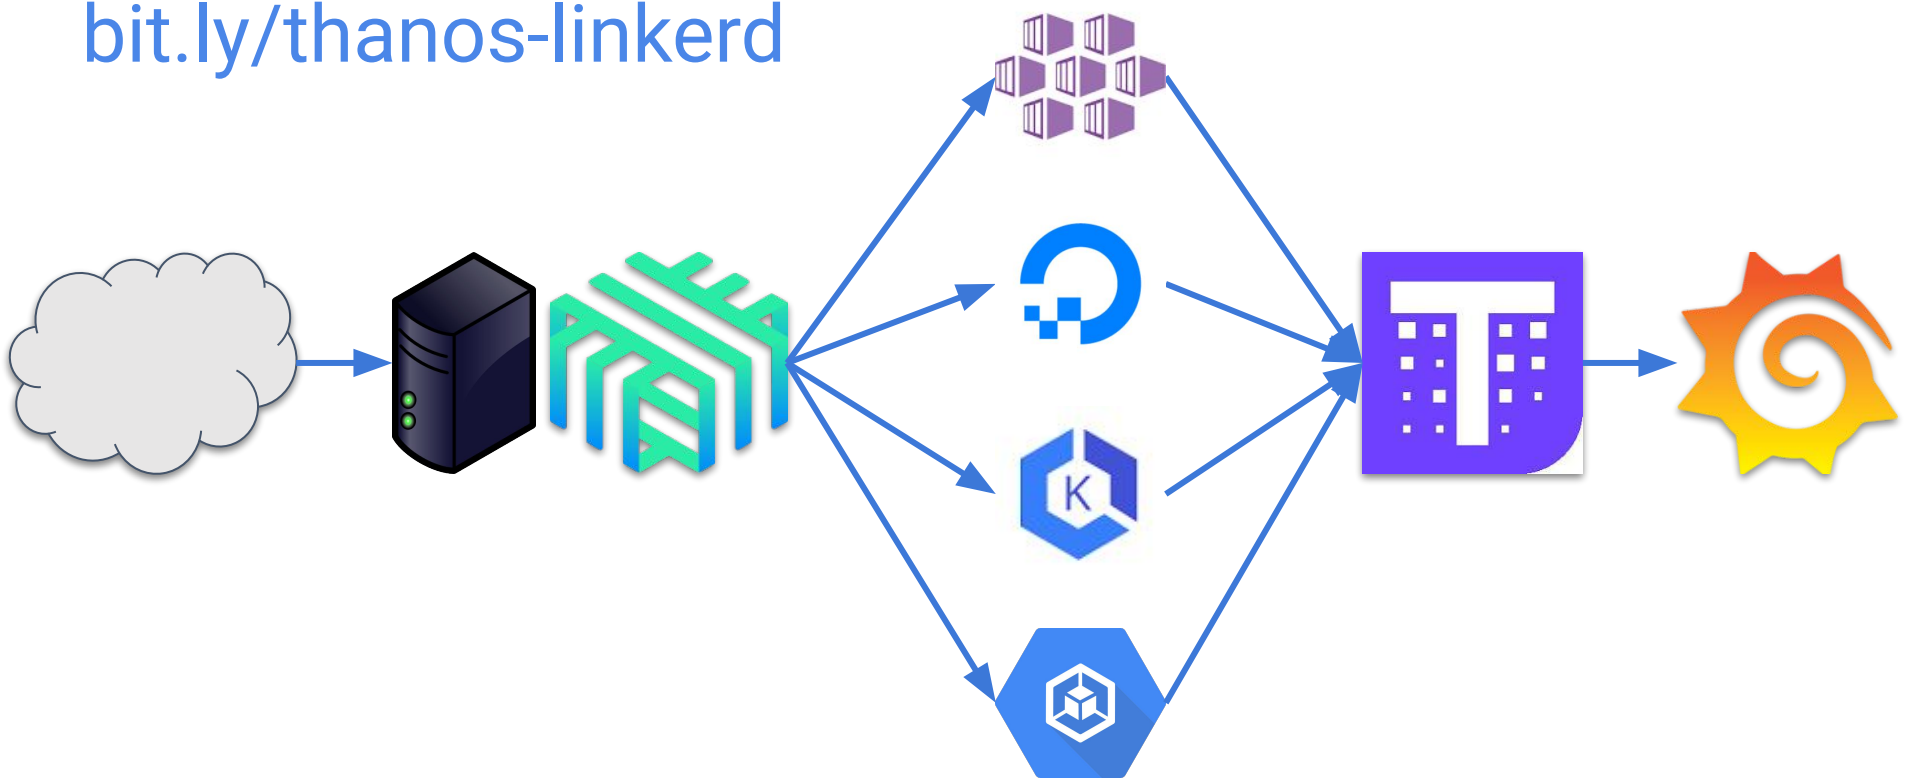

Autoscaling

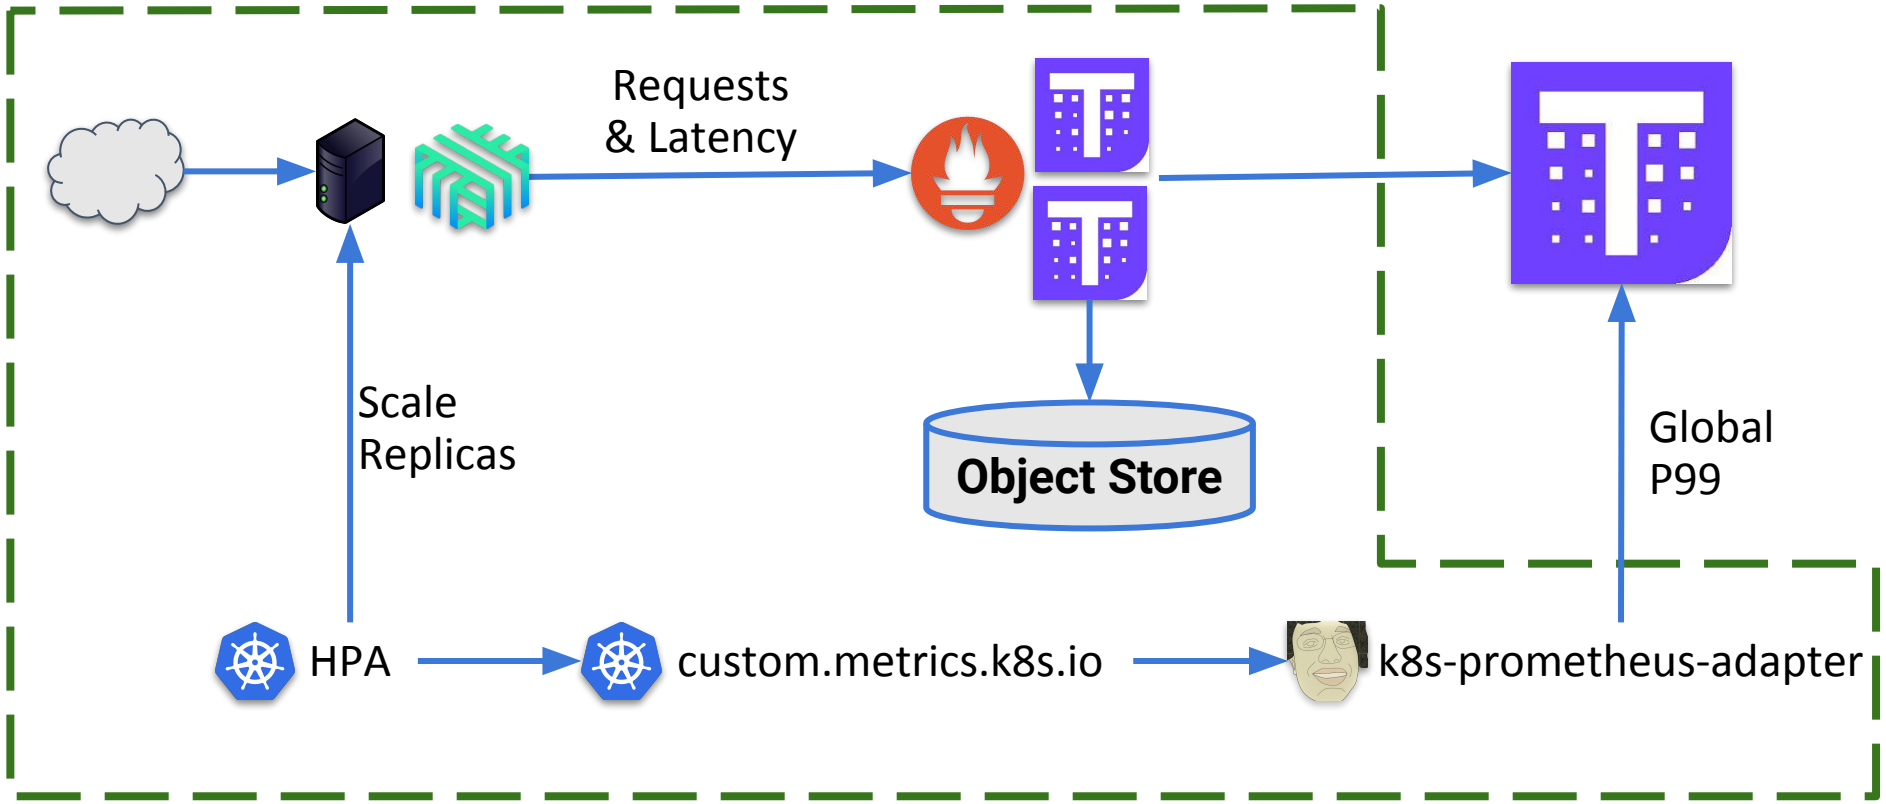

An open source *service mesh* and CNCF member project.

- 24+ months in production
- 3,000+ Slack channel members
- 10,000+ GitHub stars
- 100+ contributors

Ask
CHASE
GoDaddy
KAIROS
STRAVA
BIGCOMMERCE
Expedia
OfferUp
ebay
Cisco webex
COMCAST
planet.

bit.ly/thanos-linkerd

Thank you!

KubeCon | CloudNativeCon Europe 2019

Frederic Branczyk
@brancz @fredbrancz

Andrew Seigner
@siggy