



KubeCon



CloudNativeCon

Europe 2019

# Monitoring Service Architecture and Health with BPF

Jonathan Perry, Flowmill

# ■ Agenda

- **Demo**: visibility into **Architecture, Health and Cost**
- How that worked: **Linux + Kubernetes – “Flow monitoring”**
- Flow vs App monitoring: **Pros and Cons**
- Major challenges: **Performance & coverage** → **use eBPF**
- Building a complete system: **Collection & analysis architecture**
- Is all this really practical?: **Evaluation**
- Where next: **Adding Application monitoring** (and how)

# Hi! I'm Jonathan Perry

[jperry@flowmill.com](mailto:jperry@flowmill.com)

[www.flowmill.com](http://www.flowmill.com)

- Government: large-scale deployments
- MIT PhD: extreme monitoring systems
  - prod at Facebook
- Flowmill: CEO

# Demo application

GoogleCloudPlatform / microservices-demo

Watch ▾

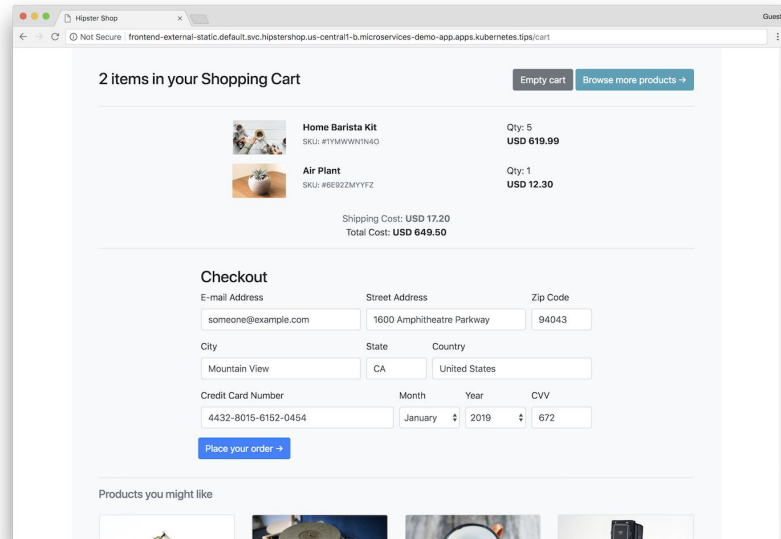
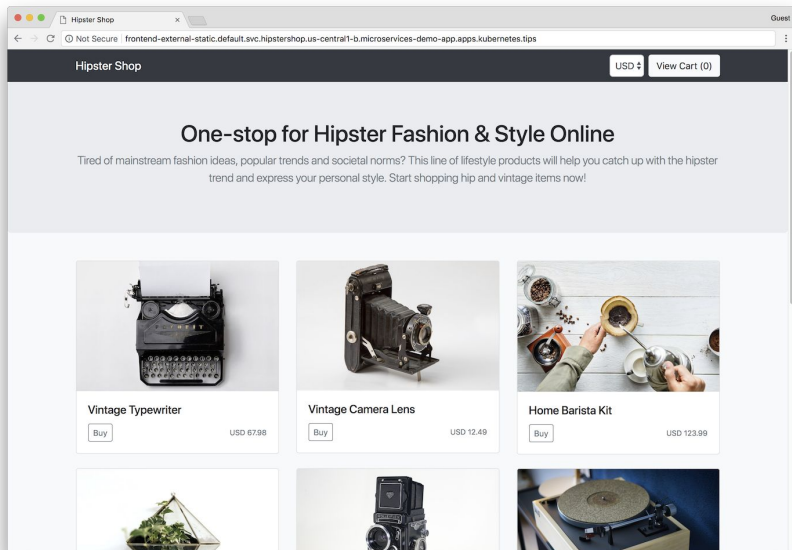
144

Unstar

4,698

Fork

596




# ■ Visibility #1: Architecture

NAME	READY	STATUS	RESTARTS	AGE
adservice-6cd6965787-lmjrl	1/1	Running	0	1d
cartservice-75f55fbc45-dv85s	1/1	Running	22	1d
checkoutservice-6848667dd7-jt44p	1/1	Running	0	1d
currencyservice-668f49f985-l42s6	1/1	Running	0	1d
emailservice-796bb9588b-cfqcp	1/1	Running	0	1d
frontend-6dcd4969b4-v2vqq	1/1	Running	0	14h
loadgenerator-54d77df7b-tgnwc	1/1	Running	0	1d
paymentservice-548657568f-p5hb7	1/1	Running	0	1d
productcatalogservice-7b94dfb45c-9gz7k	1/1	Running	0	1d
recommendationservice-5fb85f46df-hwj5s	1/1	Running	0	1d
shippingservice-5f5d75bf65-v5p7f	1/1	Running	0	1d

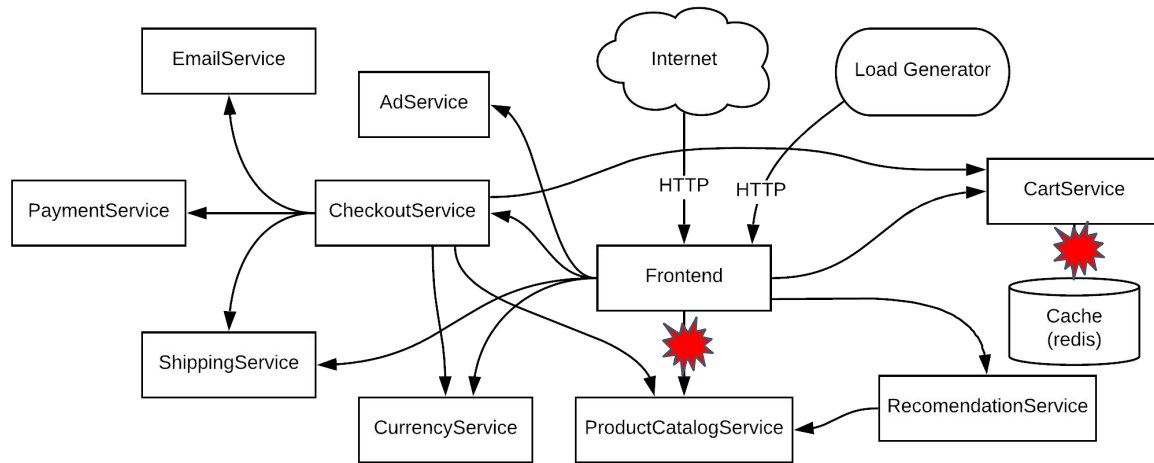
this shows the components, but how do they interact?

# ■ Visibility #1: Architecture



Just deployed, are old dependencies ok? New dependency?  
Are Dev and prod isolated?  
HA: what zones are communicating?

## ■ Visibility #2: Health

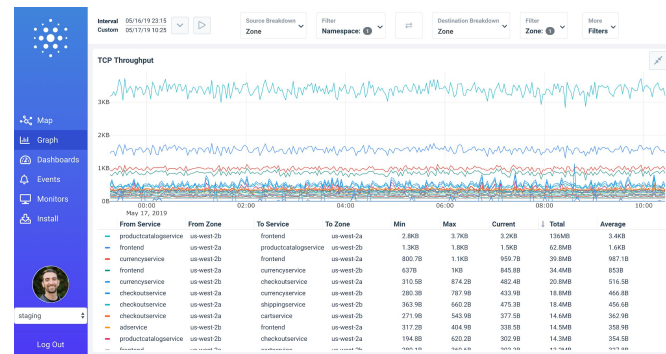


- Demo: Detecting service degradation
- Demo: Detecting security group misconfiguration

# Visibility #3: Cost

## Top service bandwidth consumption

- per Node
- across Zones → Demo
- across Regions



	From Service	From Zone	To Service	To Zone	Min	Max	Current	Total	Average
productcatalogservice	productcatalogservice	us-west-2b	frontend	us-west-2a	2.8KB	3.7KB	3.2KB	136MB	3.4KB
frontend	frontend	us-west-2a	productcatalogservice	us-west-2b	1.3KB	1.8KB	1.5KB	62.8MB	1.6KB
currencyservice	currencyservice	us-west-2b	frontend	us-west-2a	800.7B	1.1KB	959.7B	39.8MB	987.1B
frontend	frontend	us-west-2a	currencyservice	us-west-2b	637B	1KB	845.8B	34.4MB	853B
currencyservice	currencyservice	us-west-2b	checkoutservice	us-west-2a	310.5B	874.2B	482.4B	20.8MB	516.5B
checkoutservice	checkoutservice	us-west-2a	currencyservice	us-west-2b	280.3B	787.9B	433.9B	18.8MB	466.8B
checkoutservice	checkoutservice	us-west-2a	shippingservice	us-west-2b	363.9B	660.2B	475.3B	18.4MB	456.6B
checkoutservice	checkoutservice	us-west-2a	cartservice	us-west-2b	271.9B	543.9B	377.5B	14.6MB	362.9B
adservice	adservice	us-west-2b	frontend	us-west-2a	317.2B	404.9B	338.5B	14.5MB	358.9B
productcatalogservice	productcatalogservice	us-west-2b	checkoutservice	us-west-2a	194.8B	620.2B	302.9B	14.3MB	354.5B
frontend	frontend	us-west-2a	cartservice	us-west-2b	280.1B	260.6B	282.2B	12.2MB	227.8B



## ■ So far...

- Architecture
  - real-time map (when deploying)
  - high availability architecture
- Health
  - service degradation
  - security group / firewall
- Cost
  - across zones

→ with Flow Data

# How: Flow data

Linux:

Timestamp	Source	Destination	Ports	Bytes	Drops	RTT
1418530010	172.31.16.139	172.31.16.21	20641 22	4249	2	4 ms

K8s:

IP	Pod	Image	Tag	Zone
172.31.16.139	frontend	frontend-image	v1.16	us-west-1c
172.31.16.21	checkoutservice	checkout-image	v2.12a	us-west-1a

Joined:

Timestamp	Source	Destination	Ports	Bytes	Drops	RTT
1418530010	frontend frontend-image v1.16 us-west-1c	checkout checkout-image v2.12a us-west-1a	20641 22	4249	2	4 ms

# ■ Flow monitoring: Pros and cons


## Pros:

- **No code changes** – only use info from Linux+k8s
- **100% coverage** – same reason
- **Small overhead** – few, optimizable collection points
  - more on this in “Evaluation” section
- **External visibility** – observe managed services, APIs

## Cons:

- **No application-level error codes**
  - only see proxies (throughput, rtt, drops)
  - solvable – more towards end of talk

# Getting Flow Data



```
$ kubectl describe pod $POD
```

```
Name:      A
Namespace: staging
...
Status:    Running
IP:        100.101.198.137
Controlled By: ReplicaSet/A
```

```
# PID=`docker inspect -f '{{.State.Pid}}' $CONTAINER` \
nsenter -t $PID -n ss -ti
ESTAB 0 0 100.101.198.137:34940 100.65.61.118:8000
cubic wscale:9.9 rto:204 rtt:0.003/0 mss:1448 cwnd:19
ssthresh:19 bytes_acked:2525112 segs_out:15664 segs_in:15578
data_segs_out:15662 send:73365.3Mbps lastsnd:384
lastrcv:10265960 lastack:384 rcv_space:29200 minrtt:0.002
```

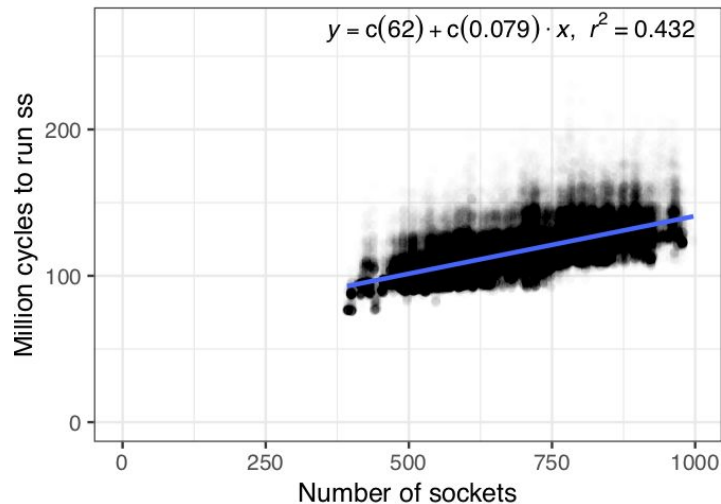
```
# conntrack -L
```

```
X tcp 6 86399 ESTABLISHED src=100.101.198.137 A
dst=100.65.61.118 sport=34940 dport=8000
B src=100.101.198.147 dst=100.101.198.137 sport=8000
dport=34940 [ASSURED] mark=0 use=1 A
```

# CLI performance & coverage

- **Performance:**

- iterates over all sockets
- built for CLI use (printfs)



- **Coverage:** Linux CLI tools are polling based



→ Misses events between polls

# ■ Enter eBPF

- Linux `bpf()` system call since 3.18
- Run code on kernel events
- Only changes, more data
- Safe: In-kernel verifier, read-only
- Fast: JIT-compiled



Unofficial BPF mascot by [Deirdré Straughan](#)

→ **100% coverage + no app changes + low overhead ftw!**

# Using eBPF

 [iovisor](#) / [bcc](#)

 Watch ▾

439

 Star

6,735

 Fork

1,130

Demo:

to run a bcc container:

```
docker run -it --rm \
  --privileged \
  -v /lib/modules:/lib/modules:ro \
  -v /usr/src:/usr/src:ro \
  -v /etc/localtime:/etc/localtime:ro \
  --workdir /usr/share/bcc/tools \
  --pid=host \
  zlim/bcc
```

<https://github.com/iovisor/bcc/blob/master/QUICKSTART.md>

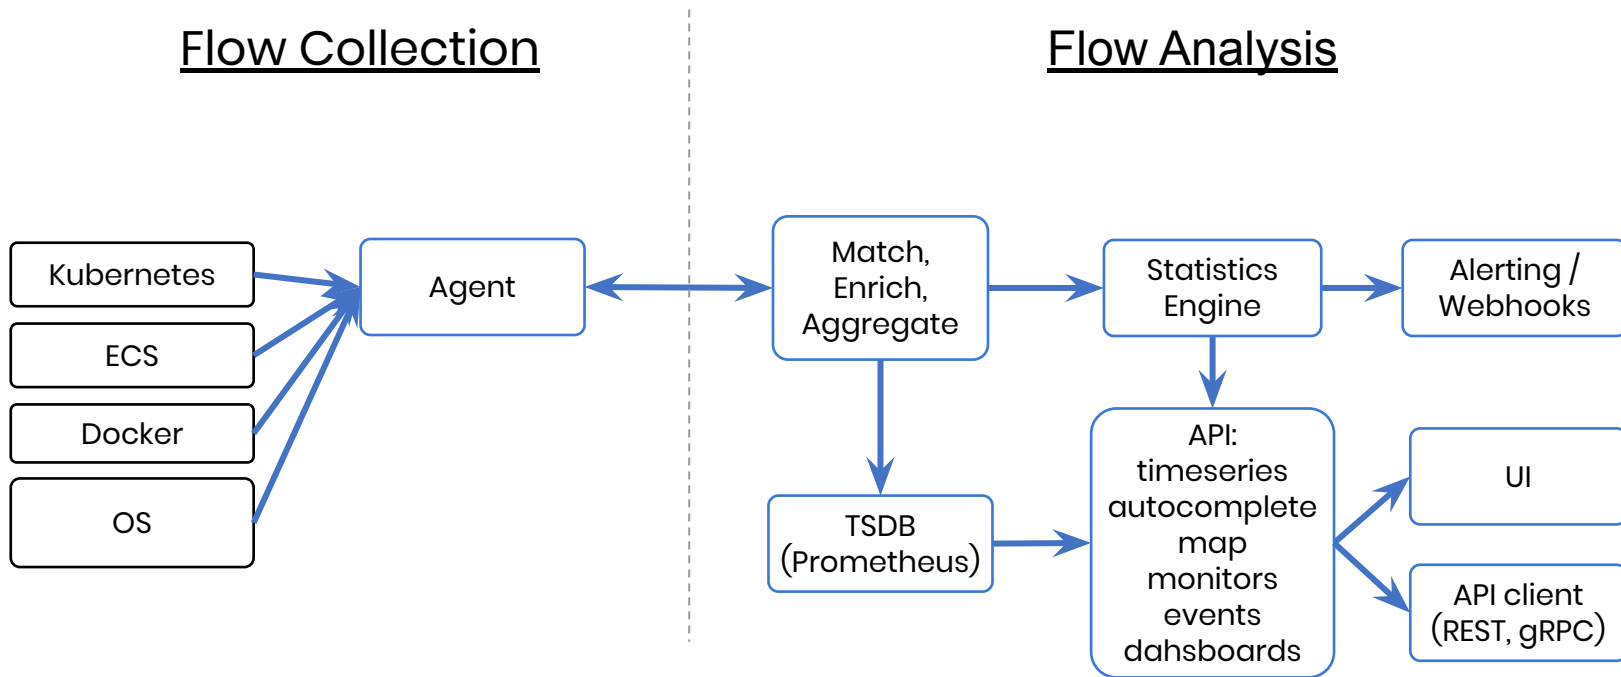
+ host pid namespace

tcptop:

- instruments `tcp_sendmsg` and `tcp_cleanup_rbuf`
- need to be careful of races:
  - # IPv4: build dict of all seen keys
  - `ipv4_throughput = defaultdict(lambda: [0, 0])`
  - for `k, v` in `ipv4_send_bytes.items()`:
  - `key = get_ipv4_session_key(k)`
  - `ipv4_throughput[key][0] = v.value`
  - `ipv4_send_bytes.clear()`

as for loop is running, kernel continues with updates, `clear()` throws those out.

# Flow monitoring: system architecture





# ■ Evaluation: CPU overhead

using *perf* and FlameGraph[1]

- To record: `perf record -a -g -e cycles -c 5000000 -- sleep 60`
- Post-process: `perf script | FlameGraph/stackcollapse-perf.pl > raw.txt`
- Analyze: `grep -E '(cleanup_module|flowmill_agent)' raw.txt |  
FlameGraph/flamegraph.pl > flame.svg`

→ observed **0.1% – 0.25% CPU overhead** across deployments

Most aggressive customer load test:

	Node	Application	TCP stack	Collector
M cycles (%)	480,000 (100%)	220,775 (46%)	27135 (5.6%)	4,120 (0.86%)

[1] [github.com/brendangregg/FlameGraph](https://github.com/brendangregg/FlameGraph)

# ■ Evaluation: Network overhead

Flow observability → monitor the flow-telemetry flows

Megabytes / second

	App throughput	Flow telemetry	%
Cluster 1	186.2	0.85	0.46%
Cluster 2	217.1	2.49	1.15%
Cluster 3	249.6	0.25	0.10%
Cluster 4 (batch)	522.0	0.16	0.031%
Cluster 5	183.0	0.02	0.013%

→ Usually < 0.5% network overhead, outliers ~1%

# ■ Evaluation: Backend QPS

Agent event counts (per second):

	TCP	UDP	NAT	process	container	DNS	Total events/s per agent
Company A	1429.2	82.0	20.8	146.5	0.014	10.5	1689.014
Company B	4017.3	89.0	-	1562.1	-	1.98	5670.38
Company C (batch)	51.0	28.8	1.05	43.8	0.55	0.5	125.7

- For a 50-node cluster, need to process 84.4k-283.5k QPS  
(~20x less for batch workloads)
- C++ analysis pipeline: hundreds of nodes w/2 second latency  
(thousands soon)

# ■ Evaluation

- CPU:
  - observed **0.1% - 0.25% CPU overhead** across deployments
  - 0.86% max load test
- Network:
  - Usually < 0.5% network overhead, outliers ~1%
- QPS:
  - ~100k QPS for 50-node cluster
  - can handle 100s of nodes with 2 second latency

# ■ Getting application error codes

- eBPF supports user probes

```
$ go tool nm /root/hello | grep 'net/http\.'
```

690a40	t	net/http.Error
64eee0	t	net/http.Get
6929e0	t	net/http.HandleFunc
6b6230	t	net/http.Handler.ServeHTTP-fm
6909e0	t	net/http.HandlerFunc.ServeHTTP
6805b0	t	net/http.Header.Add
680700	t	net/http.Header.Del
680690	t	net/http.Header.Get
680620	t	net/http.Header.Set
680750	t	net/http.Header.Write
681190	t	net/http.Header.WriteSubset
680840	t	net/http.Header.Clone

```
$ /funccount -p 31328 '/root/hello:net/http.*Header*'
Tracing 111 functions for
"/root/hello:net/http.*Header*"... Hit Ctrl-C to end.
^C
```

FUNC	COUNT
net/http.Header.Del	3
net/http.Header.sortedKeyValues	3
net/http.Header.WriteSubset	3
net/http.(*response).WriteHeader	3
net/http.extraHeader.Write	3
net/http.(*chunkWriter).writeHeader	3
net/http.(*chunkWriter).writeHeader.func1	3

```
Detaching...
```

# ■ Flow monitoring


Visibility into **Architecture, Health, and Cost**

- No code changes
- Negligible overhead
- Visibility into external dependencies
- Want application metrics (in progress)

} with eBPF

Questions? (and please reach out)

Jonathan Perry <[jperry@flowmill.com](mailto:jperry@flowmill.com)>  
[www.flowmill.com](http://www.flowmill.com)




**Backup slides**



# What about service mesh?

- Two types of data:
  - Aggregated: missing info on failure domains
  - Access logs: firehose, still need to process 100k+ events/sec



- misconfigured mesh → broken telemetry.
  - want redundant telemetry
- partial deployments & managed services